

Sourabrata Mukherjee

RESEARCH SCIENTIST · APPLIED SCIENTIST · NLP – LLMs

📍 Bengaluru, India
🌐 souro.github.io
✉️ soura1990@gmail.com

📄 in/souro
🎓 Google Scholar
🔗 souro



ABOUT ME

Research scientist and engineer working on LLM training, alignment, interpretability, and evaluation. Postdoctoral Researcher at Microsoft Research; PhD from Charles University, Prague; 6+ years shipping production ML under real-world constraints.

RESEARCH INTERESTS

- LLM Evaluation & LLM-as-Judge
- Post-training & Alignment
- Mechanistic Interpretability
- Multilingual & Cultural NLP
- Controllable Generation
- Reasoning & Agentic Systems

SKILLS & KNOWLEDGE

Training & Alignment

Full fine-tuning PEFT / LoRA
Instruction tuning DPO PPO GRPO
RLHF Reward modelling

Interpretability & Evaluation

Mechanistic interpretability Probing
LLM-as-Judge Benchmark design
Human evaluation Calibration

Engineering

Python PyTorch Hugging Face vLLM
Distributed training RAG & agents

SERVICE

Reviewer · ACL, EMNLP, NeurIPS, ICLR, COLM, INLG

LANGUAGES

Bengali ██████████
English ██████████
Hindi ██████████
Czech ██████████
German ██████████

AWARDS

- **Best Paper Award** · BLP @ EMNLP 2023
- **CU Grant Agency Award** · as PI
- **Star Employee of the Year** · Amdocs & Tricon
- **National Winner** · IBM Tech Competition

CERTIFICATION

- Generative AI with LLMs
- Deep Neural Networks & Tuning
- Machine Learning & Statistics

🎓 via Coursera

RESEARCH & SOFTWARE DEVELOPMENT

since 07/2025	Postdoctoral Researcher MICROSOFT RESEARCH 📍 BENGALURU, INDIA LLM evaluation, multilingual NLP, alignment, and controllable generation for Language AI.
from 07/2024 to 03/2025	Visiting Researcher MBZUAI 📍 ABU DHABI, UAE Cultural and cross-lingual dimensions of large language models.
from 01/2022 to 06/2022	Research Intern PANLINGUA LANGUAGE PROCESSING 📍 NEW DELHI, INDIA Low-resource machine translation for Indian languages.
from 10/2019 to 04/2025	PhD Researcher UFAL, CHARLES UNIVERSITY 📍 PRAGUE, CZECHIA Research on Text Style Transfer with neural language models.
from 12/2018 to 05/2019	Research Assistant UKP LAB, TU DARMSTADT 📍 DARMSTADT, GERMANY Context detection for scientific data-to-text generation.
from 02/2018 to 08/2018	Data Science Visiting Intern INDIAN INSTITUTE OF SCIENCE (IISc) 📍 BENGALURU, INDIA Time-series forecasting and predictive analytics.
from 06/2017 to 09/2019	Senior ML Engineer → Tech Lead TRICON INFOTECH 📍 BENGALURU, INDIA Product- and domain-specific recommendation engine at scale.
from 04/2016 to 05/2017	Senior Data Engineer AVAYA 📍 BENGALURU, INDIA Log-analysis pipeline with optimized storage and real-time monitoring.
from 05/2015 to 03/2016	Senior Analytics Engineer O9 SOLUTIONS 📍 BENGALURU, INDIA Scalable enterprise planning recommendation framework.
from 01/2014 to 04/2015	Software → Senior Software Engineer AMDOCS 📍 PUNE, INDIA Recommendation engine and search for e-commerce platforms.

EDUCATION

from 10/2019 to 04/2025	PhD in Computational Linguistics CHARLES UNIVERSITY 📍 PRAGUE, CZECHIA • Thesis: Text Style Transfer using Neural Models • Advisor: Ondřej Dušek
from 2011 to 2013	M.Tech in Computer Science & Engineering NIT DURGAPUR 📍 DURGAPUR, INDIA • GPA: 9.06 / 10 • Thesis: Real-world multi-objective optimization using evolutionary algorithms • Among the top-ranked students of the class
from 2007 to 2011	B.Tech in Computer Science & Engineering WEST BENGAL UNIVERSITY OF TECHNOLOGY 📍 KOLKATA, INDIA • GPA: 8.79 / 10 • Project: Random number generation using Monte Carlo methods • Graduated as college topper

RECENT PUBLICATIONS 🔖 View all

•	Actions Speak Louder than Words: On the Cross-Lingual Invariance of Agentic Tool Use COLM 26 · Spotlight · Top 15%
•	Building Benchmarks from the Ground Up: Community-Centered Evaluation of LLMs in Healthcare Chatbot Settings CHI 26 · Honorable Mention
•	Women, Infamous, and Exotic Beings: A Comparative Study of Honorific Usages in Wikipedia and LLMs for Bengali and Hindi EMNLP MAIN 25
•	The Geometry of LLM-as-Judge: Why Inter-LLM Consensus Is Not Human Alignment EMNLP 26 (Communicated)
•	Opening the Judge: What LLM Evaluators Compute but Do Not Report EACL 26 (Communicated)
•	Evaluation Is Training: Confidence Outruns Correctness When the Judge Is the Target EACL 26 (Communicated)