

Sourabrata Mukherjee

RESEARCH SCIENTIST · APPLIED SCIENTIST · NLP – LLMs

📍 Bengaluru, India
🌐 Personal Website
✉ Email
✉ Alternate Email

🔗 LinkedIn
🎓 Google Scholar
🐙 GitHub
🤗 Hugging Face



ABOUT ME

Research scientist and engineer working on LLM training, alignment, interpretability, and evaluation. Postdoctoral Researcher at Microsoft Research; PhD from Charles University, Prague; 6+ years shipping production ML under real-world constraints.

RESEARCH INTERESTS

- LLM Evaluation & LLM-as-Judge
- Post-training & Alignment
- Mechanistic Interpretability
- Multilingual & Cultural NLP
- Controllable Generation
- Reasoning & Agentic Systems

SKILLS & KNOWLEDGE

Training & Alignment

Full fine-tuning PEFT / LoRA
Instruction tuning DPO PPO GRPO
RLHF Reward modelling

Interpretability & Evaluation

Mechanistic interpretability Probing
LLM-as-Judge Benchmark design
Human evaluation Calibration

Engineering

Python PyTorch Hugging Face vLLM
Distributed training RAG & agents

SERVICE

Reviewer · ACL, EMNLP, NeurIPS, ICLR, COLM, INLG

LANGUAGES

Bengali	████████
English	████████
Hindi	████████
Czech	██████
German	██████

AWARDS

- **Best Paper Award** · BLP @ EMNLP 2023
- **CU Grant Agency Award** · as PI
- **Star Employee of the Year** · Amdocs & Tricon
- **National Winner** · IBM Tech Competition

CERTIFICATION

- Generative AI with LLMs
- Deep Neural Networks & Tuning
- Machine Learning & Statistics

🎓 via Coursera

RESEARCH & SOFTWARE DEVELOPMENT

since 2025

Postdoctoral Researcher

MICROSOFT RESEARCH 📍 BENGALURU, INDIA
LLM evaluation, multilingual NLP, alignment, and controllable generation for Language AI.

from 2024
to 2025

Visiting Researcher

MBZUAI 📍 ABU DHABI, UAE
Cultural and cross-lingual dimensions of large language models.

2022

Research Intern

PANLINGUA LANGUAGE PROCESSING 📍 NEW DELHI, INDIA
Low-resource machine translation for Indian languages.

from 2019
to 2025

PhD Researcher

UFAL, CHARLES UNIVERSITY 📍 PRAGUE, CZECHIA
Research on Text Style Transfer with neural language models.

from 2018
to 2019

Research Assistant

UKP LAB, TU DARMSTADT 📍 DARMSTADT, GERMANY
Context detection for scientific data-to-text generation.

2018

Data Science Visiting Intern

INDIAN INSTITUTE OF SCIENCE (IISc) 📍 BENGALURU, INDIA
Time-series forecasting and predictive analytics.

from 2016
to 2019

Senior ML Engineer → AI Tech Lead

TRICON INFOTECH 📍 BENGALURU, INDIA
Product- and domain-specific recommendation engine at scale.

2016

Senior Big Data Engineer

AVAYA 📍 BENGALURU, INDIA
Log-analysis pipeline with optimized storage and real-time monitoring.

from 2015
to 2016

Senior Analytics Engineer

O9 SOLUTIONS 📍 BENGALURU, INDIA
Scalable enterprise planning recommendation framework.

from 2014
to 2015

Software → Senior Software Engineer

AMDOCS 📍 PUNE, INDIA
Recommendation engine and search for e-commerce platforms.

EDUCATION

from 2019
to 2025

PhD in Computational Linguistics

CHARLES UNIVERSITY 📍 PRAGUE, CZECHIA
• **Thesis:** Text Style Transfer using Neural Models
• **Advisor:** Ondřej Dušek

from 2011
to 2013

M.Tech in Computer Science & Engineering

NIT DURGAPUR 📍 DURGAPUR, INDIA
• **GPA:** 9.06 / 10
• **Thesis:** Real-world multi-objective optimization using evolutionary algorithms
• Among the top-ranked students of the class

from 2007
to 2011

B.Tech in Computer Science & Engineering

WEST BENGAL UNIVERSITY OF TECHNOLOGY 📍 KOLKATA, INDIA
• **GPA:** 8.79 / 10
• **Project:** Random number generation using Monte Carlo methods
• Graduated as college topper

RECENT PUBLICATIONS

👁 View all

The Geometry of LLM-as-Judge: Why Inter-LLM Consensus Is Not Human Alignment

EMNLP 26

Actions Speak Louder than Words: Measuring Cross-Lingual Policy Retention in Tool-Using Agents

COLM 26 · Reviewers' Spotlight · Top 15%

Building Benchmarks from the Ground Up: Community-Centered Evaluation of LLMs in Healthcare Chatbot Settings

CHI 26 · Honorable Mention

Women, Infamous, and Exotic Beings: A Comparative Study of Honorific Usages in Wikipedia and LLMs for Bengali and Hindi

EMNLP 25

Opening the Judge: What LLM Evaluators Compute but Do Not Report

Communicated

Evaluation Is Training: Confidence Outruns Correctness When the Judge Is the Target

Communicated