

Sourabrata Mukherjee

RESEARCH SCIENTIST · APPLIED SCIENTIST · NLP – LLMs

📍 Bengaluru, India
🌐 souro.github.io
✉️ soura1990@gmail.com

📄 in/souro
🎓 Google Scholar
🐙 souro



ABOUT ME

Research scientist and engineer working on LLM training, alignment, interpretability, and evaluation. Postdoctoral Researcher at Microsoft Research; PhD from Charles University, Prague; 6+ years shipping production ML under real-world constraints.

RESEARCH INTERESTS

- LLM Evaluation & LLM-as-Judge
- Post-training & Alignment
- Mechanistic Interpretability
- Multilingual & Cultural NLP
- Controllable Generation
- Reasoning & Agentic Systems

SKILLS & KNOWLEDGE

Training & Alignment

Full fine-tuning PEFT / LoRA
Instruction tuning DPO PPO GRPO
RLHF Reward modelling

Interpretability & Evaluation

Mechanistic interpretability Probing
LLM-as-Judge Benchmark design
Human evaluation Calibration

Engineering

Python PyTorch Hugging Face vLLM
Distributed training RAG & agents

SERVICE

Reviewer · ACL, EMNLP, NeurIPS, ICLR, COLM, INLG

LANGUAGES

Bengali	████████
English	████████
Hindi	████████
Czech	██████
German	██████

AWARDS

- **Best Paper Award** · BLP @ EMNLP 2023
- **CU Grant Agency Award** · as PI
- **Star Employee of the Year** · Amdocs & Tricon
- **National Winner** · IBM Tech Competition

CERTIFICATION

- Generative AI with LLMs
- Deep Neural Networks & Tuning
- Machine Learning & Statistics

🎓 via Coursera

RESEARCH & SOFTWARE DEVELOPMENT

- since 07/2025
Postdoctoral Researcher
MICROSOFT RESEARCH 📍 BENGALURU, INDIA
LLM evaluation, multilingual NLP, alignment, and controllable generation for Language AI.
- from 07/2024 to 03/2025
Visiting Researcher
MBZUAI 📍 ABU DHABI, UAE
Cultural and cross-lingual dimensions of large language models.
- from 01/2022 to 06/2022
Research Intern
PANLINGUA LANGUAGE PROCESSING 📍 NEW DELHI, INDIA
Low-resource machine translation for Indian languages.
- from 10/2019 to 04/2025
PhD Researcher
UFAL, CHARLES UNIVERSITY 📍 PRAGUE, CZECHIA
Research on Text Style Transfer with neural language models.
- from 12/2018 to 05/2019
Research Assistant
UKP LAB, TU DARMSTADT 📍 DARMSTADT, GERMANY
Context detection for scientific data-to-text generation.
- from 02/2018 to 08/2018
Data Science Visiting Intern
INDIAN INSTITUTE OF SCIENCE (IISc) 📍 BENGALURU, INDIA
Time-series forecasting and predictive analytics.
- from 06/2017 to 09/2019
Senior ML Engineer → Tech Lead
TRICON INFOTECH 📍 BENGALURU, INDIA
Product- and domain-specific recommendation engine at scale.
- from 04/2016 to 05/2017
Senior Data Engineer
AVAYA 📍 BENGALURU, INDIA
Log-analysis pipeline with optimized storage and real-time monitoring.
- from 05/2015 to 03/2016
Senior Analytics Engineer
O9 SOLUTIONS 📍 BENGALURU, INDIA
Scalable enterprise planning recommendation framework.
- from 01/2014 to 04/2015
Software → Senior Software Engineer
AMDOCS 📍 PUNE, INDIA
Recommendation engine and search for e-commerce platforms.

EDUCATION

- from 10/2019 to 04/2025
PhD in Computational Linguistics
CHARLES UNIVERSITY 📍 PRAGUE, CZECHIA
• **Thesis:** Text Style Transfer using Neural Models
• **Advisor:** Ondřej Dušek
- from 2011 to 2013
M.Tech in Computer Science & Engineering
NIT DURGAPUR 📍 DURGAPUR, INDIA
• **GPA:** 9.06 / 10
• **Thesis:** Real-world multi-objective optimization using evolutionary algorithms
• Among the top-ranked students of the class
- from 2007 to 2011
B.Tech in Computer Science & Engineering
WEST BENGAL UNIVERSITY OF TECHNOLOGY 📍 KOLKATA, INDIA
• **GPA:** 8.79 / 10
• **Project:** Random number generation using Monte Carlo methods
• Graduated as college topper

SELECTED PUBLICATIONS 🔖 View all

- **Building Benchmarks from the Ground Up: Community-Centered Evaluation of LLMs in Healthcare Chatbot Settings**
CHI 26 · Honorable Mention
- **Women, Infamous, and Exotic Beings: A Comparative Study of Honorific Usages in Wikipedia and LLMs for Bengali and Hindi**
EMNLP MAIN 25
- **Actions Speak Louder than Words: On the Cross-Lingual Invariance of Agentic Tool Use**
COLM 26 · Spotlight · Top 15%
- **The Geometry of LLM-as-Judge: Why Inter-LLM Consensus Is Not Human Alignment**
EMNLP 26 (Communicated)
- **Opening the Judge: What LLM Evaluators Compute but Do Not Report**
EACL 26 (Communicated)
- **Evaluation Is Training: Confidence Outruns Correctness When the Judge Is the Target**
EACL 26 (Communicated)